

Supplementary

# Outline

- Tool
  - Chisel
  - Rocket-chip
  - Chipyard
  - Riscv-tests
- HW5

# Outline

- Tool
  - Chisel
  - Rocket-chip
  - Chipyard
  - Riscv-tests
- HW5

# Chisel

- Constructing hardware in a Scala embedded language
- Open source hardware description language to describe circuits at RTL
- It can be converted to verilog for synthesis
- Not exactly HLS cause there's no EDA tool for it
- Developed by Berkeley

```
class Add extends Module {  
  val io = IO(new Bundle {  
    val a = Input(UInt(8.W))  
    val b = Input(UInt(8.W))  
    val y = Output(UInt(8.W))  
  })  
  
  io.y := io.a + io.b  
}
```

A 32-bit register with a reset value of 0:


```
val reg = RegInit(0.U(32.W))
```

A multiplexer is part of the Chisel library:

```
val result = Mux(sel, a, b)
```

# Chisel example

- Example of FIRfilter



While Chisel provides similar base primitives as synthesizable Verilog, and *could* be used as such:

The above can be converted to Verilog using `ChiselStage`:

```
import chisel3.stage.{ChiselStage, ChiselGeneratorAnnotation}

(new chisel3.stage.ChiselStage).execute(
  Array("-X", "verilog"),
  Seq(ChiselGeneratorAnnotation(() => new FirFilter(8, Seq(1.U, 1.U, 1.U)))))
```

Alternatively, you may generate some Verilog directly for inspection:

```
val verilogString = (new chisel3.stage.ChiselStage).emitVerilog(new FirFilter(8, Seq(0.U, 1.U)))
println(verilogString)
```

Generate Verilog file

// 3-point moving average implemented in the style of a FIR filter

```
class MovingAverage3(bitWidth: Int) extends Module {
  val io = IO(new Bundle {
    val in = Input(UInt(bitWidth.W))
    val out = Output(UInt(bitWidth.W))
  })
```

i/o specification

```
  val z1 = RegNext(io.in)
  val z2 = RegNext(z1)
```

Sequential part

```
  io.out := (io.in * 1.U) + (z1 * 1.U) + (z2 * 1.U)
}
```

Combinational part

# Rocket-chip

- Rocket chip generator is an SoC generator developed at Berkeley and now supported by SiFive
- High extensibility with easy configuration

```
class RocketConfig extends Config(  
  new chipyard.iobinders.WithUARTAdapter ++  
  new chipyard.iobinders.WithTieOffInterrupts ++  
  new chipyard.iobinders.WithBlackBoxSimMem ++  
  new chipyard.iobinders.WithTieOffDebug ++  
  new chipyard.iobinders.WithSimSerial ++  
  new testchipip.WithTSI ++  
  new chipyard.config.WithBootROM ++  
  new chipyard.config.WithUART ++  
  new chipyard.config.WithL2TLBs(1024) ++  
  new freechips.rocketchip.subsystem.WithNoMMIOPort ++  
  new freechips.rocketchip.subsystem.WithNoSlavePort ++  
  new freechips.rocketchip.subsystem.WithInclusiveCache ++  
  new freechips.rocketchip.subsystem.WithNExtTopInterrupts(0) ++  
  new freechips.rocketchip.subsystem.WithNBigCores(1) ++  
  new freechips.rocketchip.subsystem.WithCoherentBusTopology ++  
  new freechips.rocketchip.system.BaseConfig)
```

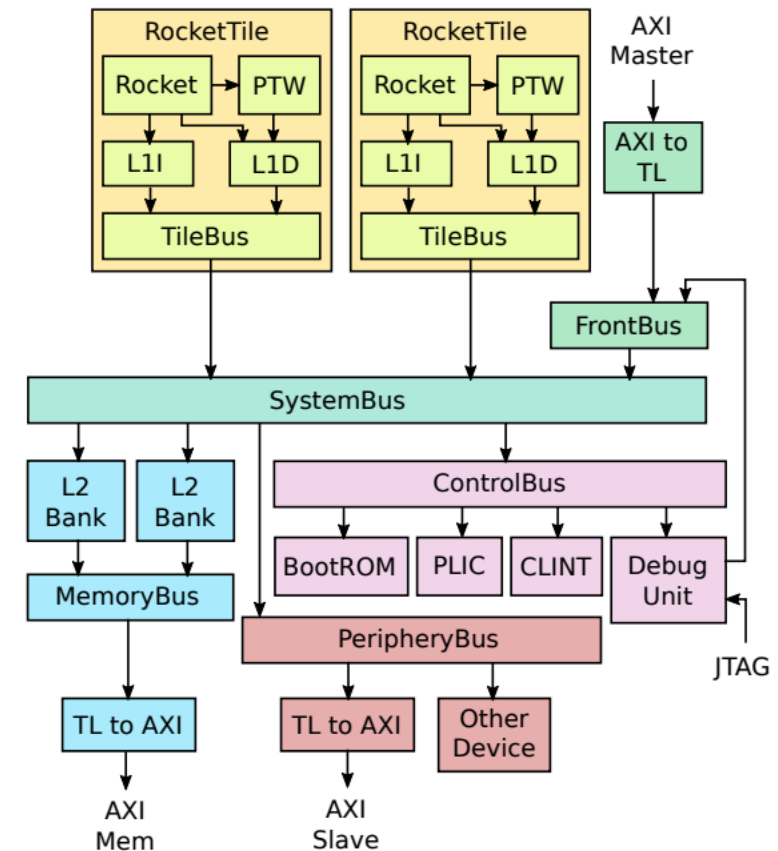
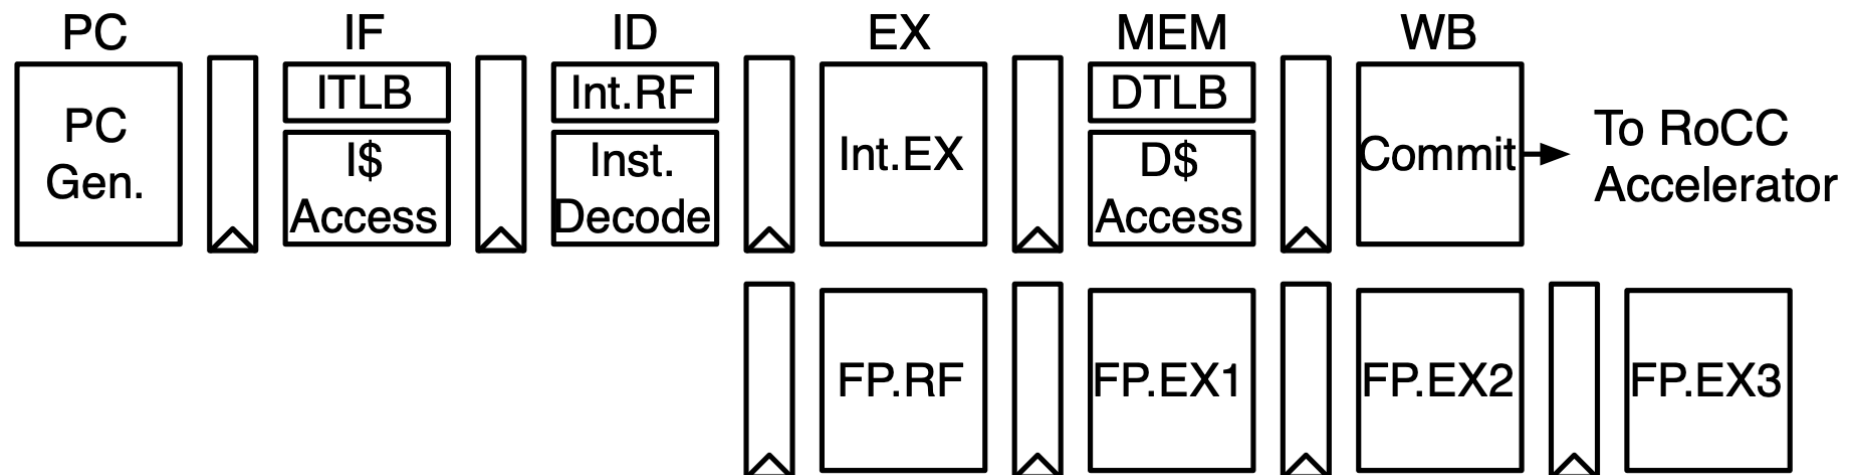


Figure 2: A generic Rocket Chip instance

# Rocket-core

- Default rocket core is a 5-stage single-issue in-order pipeline
- 64-entry BTB, 256-entry BHT, 2-entry RAS
- MMU supports page-based virtual memory
- IEEE 754-2008-compliant FPU

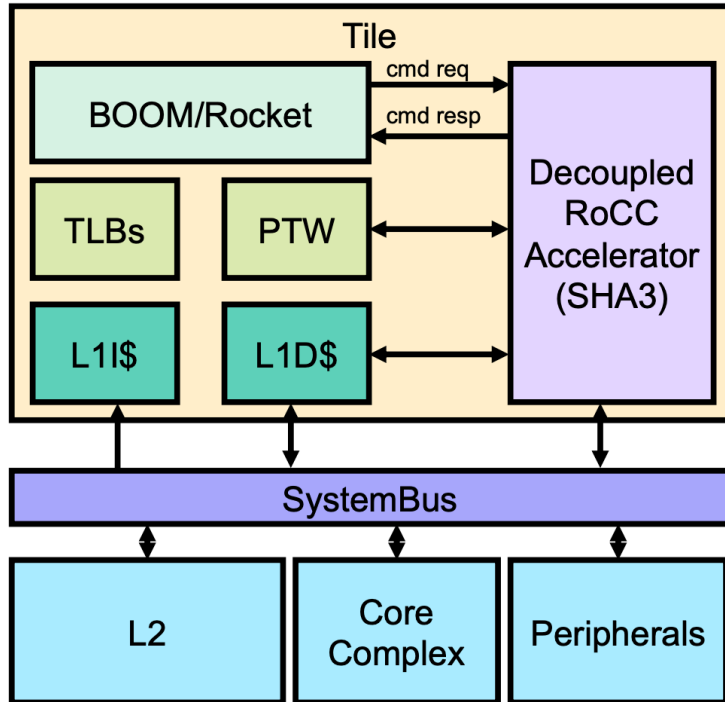


# Chipyard

- Integrate lots of packages, including rocket-chip, firrtl, sifive-cache, gemmini, nvdla, sha3, etc.



# Add a sha3 coprocessor



Execute custom RISC-V instructions for a custom extension

31	25	24	20	19	15	14	13	12	11	7	6	0
funct	rs2	rs1	xd	xs1	xs2	rd	opcode					
7	5	5	1	1	1	5	7					

```
class TutorialRocketConfig extends Config(
  new WithTop ++
  new WithBootROM ++
  new freechips.rocketchip.subsystem.WithInclusiveCache ++
  new freechips.rocketchip.subsystem.WithNBigCores(1) ++
  new freechips.rocketchip.system.BaseConfig)
```

Modify and add

```
class Sha3RocketConfig extends Config(
  new WithTop ++
  new WithBootROM ++
  new freechips.rocketchip.subsystem.WithInclusiveCache ++
  new sha3.WithSha3Accel ++
  new freechips.rocketchip.subsystem.WithNBigCores(1) ++
  new freechips.rocketchip.system.BaseConfig)
```

Expanded C macro

```
asm volatile ("custom2 x0, %[msg_addr], %[hash_addr], 0"
  :: [msg_addr] "r" (&input), [hash_addr] "r" (&output));
```

# Incorporating verilog blocks

```
generators/yourproject/  
  build.sbt  
  src/main/  
    scala/  
      resources/  
        vsrc/  
          YourFile.v
```

## Chisel BlackBox Definition

```
class GCDMMIOBlackBox(val w: Int) extends BlackBox(Map("WIDTH" -> IntParam(w))) with HasBlackBoxIO  
  with HasGCDIO  
{  
  addResource("/vsrc/GCDMMIOBlackBox.v")  
}
```

```
module GCDMMIOBlackBox  
  #(parameter WIDTH)  
  (  
    input          clock,  
    input          reset,  
    output         input_ready,  
    input          input_valid,  
    input [WIDTH-1:0] x,  
    input [WIDTH-1:0] y,  
    input          output_ready,  
    output         output_valid,  
    output reg [WIDTH-1:0] gcd,  
    output         busy  
  );
```

```
class GCDAXI4BlackBoxRocketConfig extends Config(  
  new chipyard.iobinders.WithUARTAdapter ++  
  new chipyard.iobinders.WithTieOffInterrupts ++  
  new chipyard.iobinders.WithBlackBoxSimMem ++  
  new chipyard.iobinders.WithTiedOffDebug ++  
  new chipyard.iobinders.WithSimSerial ++  
  new testchipip.WithTSI ++  
  new chipyard.config.WithUART ++  
  new chipyard.config.WithBootROM ++  
  new chipyard.config.WithL2TLBs(1024) ++  
  new chipyard.example.WithGCD(useAXI4=true, useBlackBox=true) ++           // Use GCD blackbox  
  new freechips.rocketchip.subsystem.WithNoMMIOPort ++  
  new freechips.rocketchip.subsystem.WithNoSlavePort ++  
  new freechips.rocketchip.subsystem.WithInclusiveCache ++  
  new freechips.rocketchip.subsystem.WithNExtTopInterrupts(0) ++  
  new freechips.rocketchip.subsystem.WithNBigCores(1) ++  
  new freechips.rocketchip.subsystem.WithCoherentBusTopology ++  
  new freechips.rocketchip.system.BaseConfig)
```

# Riscv-tests

- Contain ISA test (privileged and non-privileged) and some small benchmark tests
- Small benchmark tests
  - Qsort
  - Rsort
  - [Dhrystone](#)
  - Median
  - Tower
  - Mm
  - Multiply
  - Mt-matmul
  - ...

# Outline

- Tool
  - Chisel
  - Rocket-chip
  - Chipyard
  - Riscv-tests
- HW5

# Memory hierarchy

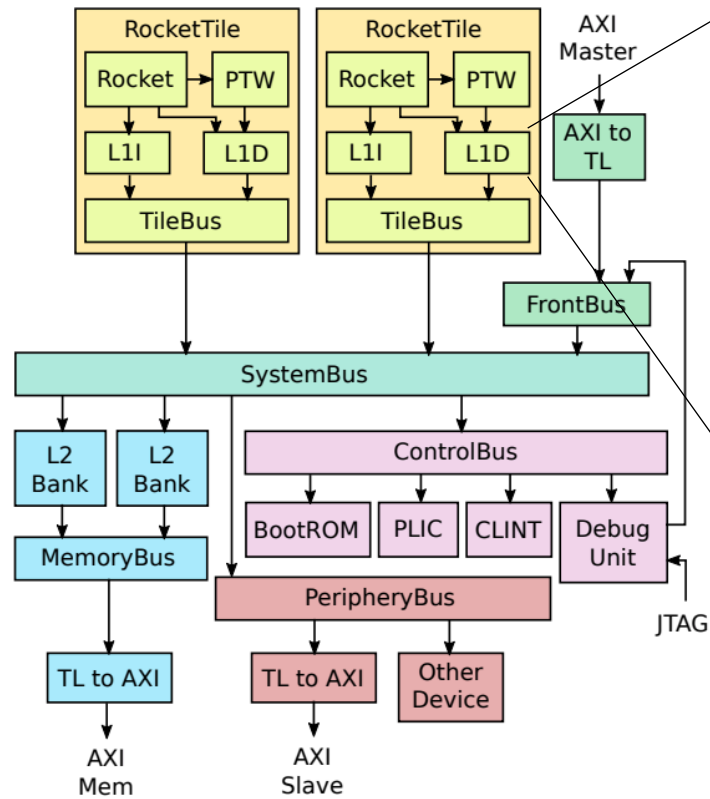
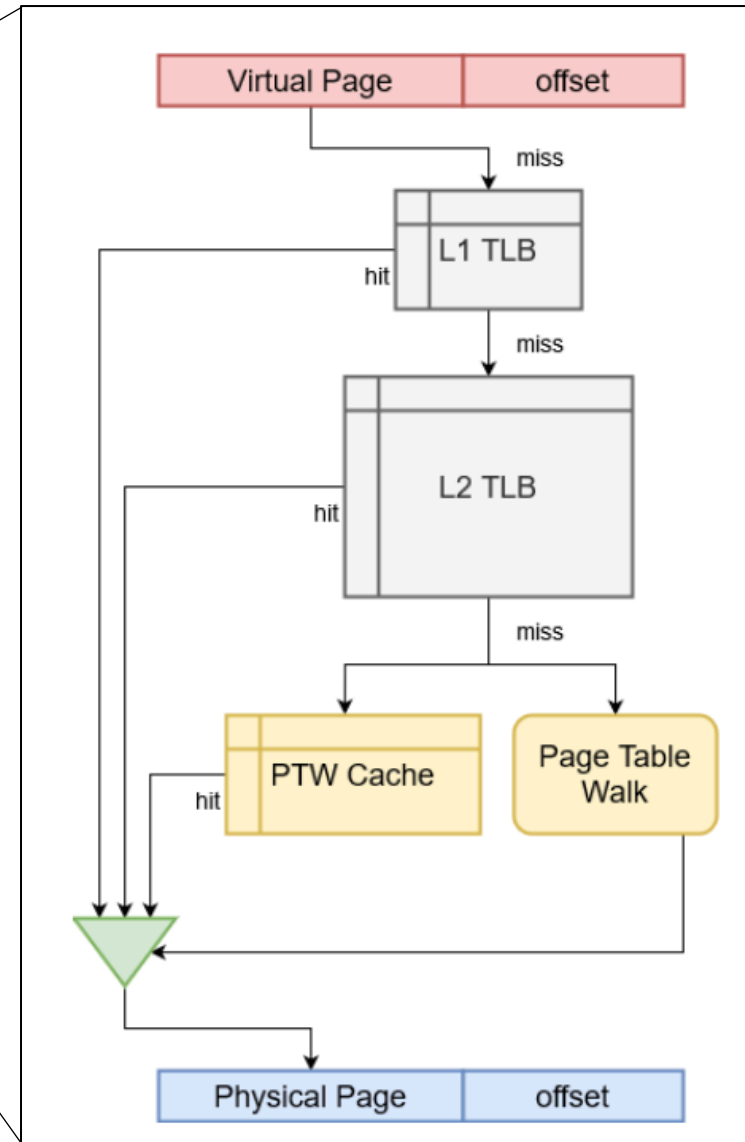


Figure 2: A generic Rocket Chip instance



# Configurable cache SoC

/root/emulator/system/test.scala

```
class Config1 extends Config(  
  // new freechips.rocketchip.subsystem.WithInclusiveCache(outerLatencyCycles = 40, capacityKB = 2,  
  // nBanks = 1, nWays = 1) ++ // L2 Cache  
  new chipyard.config.WithL2TLBs(0) ++  
  new WithNTestCores(n_core = 1,  
    L1_Dcache_Set = 16, L1_Dcache_Way = 1,  
    L1_Icache_Set = 16, L1_Icache_Way = 1,  
    L1_Dcache_replacement = "random", // "lru", "plru"  
    L1_Dcache_TLB_Set = 1, L1_Dcache_TLB_Way = 4,  
    L1_Icache_TLB_Set = 1, L1_Icache_TLB_Way = 4,  
    L2_TLB_Set = 0, L2_TLB_Way = 4,  
    enable_vm = false  
  ) ++  
  new TestBaseConfig()  
)
```

```
class Config2 extends Config(  
  // new freechips.rocketchip.subsystem.WithInclusiveCache(outerLatencyCycles = 40, capacityKB = 2,  
  // nBanks = 1, nWays = 1) ++ // L2 Cache  
  new chipyard.config.WithL2TLBs(0) ++  
  new WithNTestCores(n_core = 1,  
    L1_Dcache_Set = 8, L1_Dcache_Way = 2,  
    L1_Icache_Set = 16, L1_Icache_Way = 1,  
    L1_Dcache_replacement = "random", // "lru", "plru"  
    L1_Dcache_TLB_Set = 1, L1_Dcache_TLB_Way = 4,  
    L1_Icache_TLB_Set = 1, L1_Icache_TLB_Way = 4,  
    L2_TLB_Set = 0, L2_TLB_Way = 4,  
    enable_vm = false  
  ) ++  
  new TestBaseConfig()  
)
```

# Some constraints

- Block size (L1\$,L2\$): 64 bytes
  - Not configurable, related to bus width
- L2 inclusive cache at least 2 ways
  - $nBanks * nWays * nSets * 64 = \text{capacity}$
- L1, L2 TLB at least 4 ways
  - May not be used in this homework because of the small size of benchmarks
- $L2\_TLB\_Set * L2\_TLB\_Way = \text{chipyard.config.WithL2TLBs(...)}$ 
  - Should specify both explicitly
- Replacement policy only affects L1D\$
- Way, Set, Bank should be 2, 4, 8, 16, ...

```
class Config1 extends Config(
  // new freechips.rocketchip.subsystem.WithInclusiveCache(outerLatencyCycles = 40, capacityKB = 2,
  // nBanks = 1, nWays = 1) ++ // L2 Cache
  new chipyard.config.WithL2TLBs(0) ++
  new WithTestCores(n_core = 1,
    L1_Dcache_Set = 16, L1_Dcache_Way = 1,
    L1_Icache_Set = 16, L1_Icache_Way = 1,
    L1_Dcache_replacement = "random", // "lru", "plru"
    L1_Dcache_TLB_Set = 1, L1_Dcache_TLB_Way = 4,
    L1_Icache_TLB_Set = 1, L1_Icache_TLB_Way = 4,
    L2_TLB_Set = 0, L2_TLB_Way = 4,
    enable_vm = false
  ) ++
  new TestBaseConfig()
)
```

# Some other things

- If L2\$ is not enabled, when L1 misses, the outer access latency is determined by bus width and L1 cache blocks, and hence, it is determined by the number of beats which is 8 cycles.
- If L2\$ is enabled, L1 to L2 is still 8 cycles (perhaps), and L2 to outer memory is 40 cycles (by default).
- So when L2 is enabled, the cycle count increases but with more realistic modeling.



# Part1

- You can collect part1 info by running ./test.sh

```
echo "dhrystone:"
for i in {1..13};
do
    echo $i
    ./Config$i benchmarks/dhrystone.riscv
done

echo ""
echo "median:"
for i in {1..13};
do
    echo $i
    ./Config$i benchmarks/median.riscv
done

echo ""
echo "multiply:"
for i in {1..13};
do
    echo $i
    ./Config$i benchmarks/multiply.riscv
done
```

# Part1

- As for 1, 2, 4 cores setting, you should change the number of the available cores in /root/emulator/benchmarks/common/crt.S and recompile the program

```
# get core id
csrr a0, mhartid
# for now, assume only 1 core
li a1, 4
# 1:bgeu a0, a1, 1b
```

→ 1, 2, 4

```
root@e4dbeab9eb48:~/emulator/benchmarks# make
```

Recompile the program (It will recompile not only mt-matmul, but the whole benchmarks)

```
root@e4dbeab9eb48:~/emulator# ./Config17 benchmarks/mt-matmul.riscv → 1
```

```
root@e4dbeab9eb48:~/emulator# ./Config19 benchmarks/mt-matmul.riscv → 2
```

```
root@e4dbeab9eb48:~/emulator# ./Config20 benchmarks/mt-matmul.riscv → 4
```

# Part2

L1 I\$ D\$ available settings:  
(set, way)  
(16, 1)  
(8, 2)  
(4, 4)  
(2, 8)

L1 D\$ replacement  
available settings:  
random  
lru  
plru

Some combinations may not work !!

- Configure the HW5.scala file
- To keep it simple, only L1 I\$ and L1 D\$ will be configured

/root/emulator/system/HW5.scala

Can be modified

```
class HW5Config extends Config(  
  // new freechips.rocketchip.subsystem.WithInclusiveCache(outerLatencyCycles = 40, capacityKB = 4, // disabled  
  // nBanks = 4, nWays = 2) ++  
  // new chipyard.config.WithL2TLBs(0) ++  
  new WithNTTestCores(n_core = 4,  
    L1_Dcache_Set = 16, L1_Dcache_Way = 1,  
    L1_Icache_Set = 16, L1_Icache_Way = 1,  
    L1_Dcache_replacement = "random",  
    L1_Dcache_TLB_Set = 1, L1_Dcache_TLB_Way = 4, // disabled  
    L1_Icache_TLB_Set = 1, L1_Icache_TLB_Way = 4, // disabled  
    L2_TLB_Set = 0, L2_TLB_Way = 4, // disabled  
    enable_vm = false  
  ) ++  
  new freechips.rocketchip.system.TestBaseConfig()  
)
```

# Change the testdata to 64x64 in mt\_matmul.c

/root/emulator/benchmarks/mt-matmul/mt\_matmul.c

```
//-----  
// Input/Reference Data  
  
#include "dataset2.h"  
  
//-----  
// Basic Utilities and Multi-thread Support  
  
#include "util.h"
```

dataset.h → 32x32

dataset2.h → 64x64

# Also change number of cores to 4

```
# get core id
csrr a0, mhartid
# for now, assume only 1 core
li a1, 4
# 1:bgeu a0, a1, 1b
```

4

```
root@e4dbeab9eb48:~/emulator/benchmarks# make
```

compile

# Compile the target configuration

```
class HW5Config extends Config(  
  // new frechips.rocketchip.subsys  
  // nBanks = 4, nWays = 2) ++  
  // new chipyard.config.WithL2TLBs(  
    new WithNTestCores(n_core = 4,
```

Class name ↑

`rm -rf generated-src` (if you build the emulator before)

```
root@e4dbeab9eb48:~/# cd /root/emulator
```

```
root@e4dbeab9eb48:~/emulator# make -j8 CONFIG=frechips.rocketchip.system.HW5Config
```

- We can see the address mapping

## Generated Address Map

0	-	1000	ARWX	debug-controller@0
3000	-	4000	ARWX	error-device@3000
10000	-	20000	R X	rom@10000
2000000	-	2010000	ARW	clint@2000000
2010000	-	2011000	ARW	cache-controller@2010000
c000000	-	10000000	ARW	interrupt-controller@c000000
60000000	-	80000000	RWX	mmio-port-axi4@60000000
80000000	-	c0000000	ARWXC	memory@80000000

Run the test

Generated binary ↗

Test program ↑

```
root@e4dbeab9eb48:~/emulator# ./emulator-frechips.rocketchip.system-frechips.rocketchip.system.HW5Config benchmark/mt-vvadd.riscv
```

```
This emulator compiled with JTAG Remote Bitbang client. To enable, use +jtag_rbb_enable=1.  
Listening on port 35255
```

# mt-matmul

/root/emulator/benchmarks/mt-matmul/

matmul.c

```
#include "dataset.h"
#include "util.h"
#include <stddef.h>

#pragma GCC optimize ("unroll-loops")

void matmul(const size_t coreid, const size_t ncores, const size_t lda,
const data_t A[], const data_t B[], data_t C[])
{
    size_t i, j, k;
    size_t block = lda / ncores;
    size_t start = block * coreid;

    for (i = 0; i < lda; i++) {
        for (j = start; j < (start+block); j++) {
            data_t sum = 0;
            for (k = 0; k < lda; k++)
                sum += A[j*lda + k] * B[k*lda + i];
            C[i + j*lda] = sum;
        }
    }
}
```

You may use inline assembly if you think compiler is not doing well  
<https://gcc.gnu.org/onlinedocs/gcc/RISC-V-Options.html>

mt-matmul.c

```
//-----
// Main
//
// all threads start executing thread_entry(). Use their "coreid" to
// differentiate between threads (each thread is running on a separate core).
void thread_entry(int cid, int nc)
{
    static data_t results_data[ARRAY_SIZE];

    stats(matmul(cid, nc, DIM_SIZE, input1_data, input2_data, results_data); barrier(nc), DIM_SIZE/DIM_SIZE/DIM_SIZE);

    int res = verify(ARRAY_SIZE, results_data, verify_data);

    exit(res);
}
```

You may also want to use barrier synchronization in your matmul.c when needed.

# Evaluation

- We will put your matmul.c in matmul/ and **make** under benchmark/
- We will also **make CONFIG=freechips.rocketchip.system.HW5Config** under emulator/
- We will score based on the cycle count
- Correctness is a must



# Default mt-matmul (64x64) status

5~7 minutes

```
This emulator compiled with JTAG Remote Bitbang client. To enable, use +jtag_rbb_enable=1.  
Listening on port 33743
```

```
matmul(cid, nc, 64, input1_data, input2_data, results_data); barrier(nc) 3544288 cycles, 13.5 cycles/iter, 8.4  
CPI
```

```
This emulator compiled with JTAG Remote Bitbang client. To enable, use +jtag_rbb_enable=1.  
Listening on port 39890
```

# You may want to test using spike before cycle accurate testing

No failed  
message

```
root@884a6d6d04e9:~/emulator# spike -p4 --ic=16:1:64 --dc=16:1:64 benchmarks/mt-matmul.riscv
matmul(cid, nc, 64, input1_data, input2_data, results_data); barmatmul(cid, nc, 64, input1_data, input2_data, results_data); barmatmul(cid, nc, 64, input1_data, input2_data, results_data); barmatmul(cid, nc, 64, input1_data, input2_data, results_data); barrier(nc): 369418 cycles, 1.4 cycles/iter, 0.9 CPI
rier(nc): 369418 cycles, 1.4 cycles/iter, 0.9 CPI
rier(nc): 369418 cycles, 1.4 cycles/iter, 0.9 CPI
rier(nc): 368249 cycles, 1.4 cycles/iter, 0.9 CPI
D$ Bytes Read:      2495309
D$ Bytes Written:    21232
D$ Read Accesses:    579734
D$ Write Accesses:   5882
D$ Read Misses:      301869
D$ Write Misses:     4164
D$ Writebacks:       4190
D$ Miss Rate:        52.258%
I$ Bytes Read:       5041272
I$ Bytes Written:     0
I$ Read Accesses:    1605040
I$ Write Accesses:    0
I$ Read Misses:      82
I$ Write Misses:     0
I$ Writebacks:       0
I$ Miss Rate:        0.005%
```

# Submission format and deadline

## Submission

- Zip and upload your file to ceiba in the following format:

```
r09922028/          <-- zip this folder
  |-- matmul.c       // Matrix multiplication program
  |-- HW5.scala      // Cache configuration
  '-- HW5.pdf        // Including handwritten part, programming part 1 answers,
                     // report on part 2 multi-core matrix-multiplication,
                     // and bonus part if any
```

- Deadline: 1/11, 21:00
- **No late submission!!!**
- If there's any question, please send email to [ntuca2020@gmail.com](mailto:ntuca2020@gmail.com).
- TA hour for this homework: Wed 15:00 17:00 and Thur 15:00 17:00

# Reference

- Chisel
  - [https://en.wikipedia.org/wiki/Chisel\\_\(programming\\_language\)](https://en.wikipedia.org/wiki/Chisel_(programming_language))
  - <https://www.chisel-lang.org/>
- Rocket-chip
  - <https://github.com/chipsalliance/rocket-chip>
  - <https://chipyard.readthedocs.io/en/latest/Generators/Rocket-Chip.html>
  - <https://riscv.org/wp-content/uploads/2015/01/riscv-rocket-chip-generator-workshop-jan2015.pdf>
- Chipyard
  - <https://chipyard.readthedocs.io/en/latest/Chipyard-Basics/index.html>
  - [https://fires.im/micro19-slides-pdf/03\\_building\\_custom\\_soc.pdf](https://fires.im/micro19-slides-pdf/03_building_custom_soc.pdf)
- Dhrystone: <https://en.wikipedia.org/wiki/Dhrystone>